

Congratulations! You passed!

Grade received **100%** Latest Submission Grade 100% To pass 80% or higher


[Go to next item](#)

1.

1 / 1 point

The Lunar Lander is a continuous state Markov Decision Process (MDP) because:

- ☐ The state-action value $Q(s, a)$ function outputs continuous valued numbers
- ☐ The reward contains numbers that are continuous valued
- ☒ The state contains numbers such as position and velocity that are continuous valued.
- ☐ The state has multiple numbers rather than only a single number (such as position in the x -direction)

 **Correct**
That's right!

2.

1 / 1 point

In the learning algorithm described in the videos, we repeatedly create an artificial training set to which we apply supervised learning where the input $x = (s, a)$ and the target, constructed using Bellman's equations, is $y = \underline{\hspace{1cm}}$?

- ☐ $y = R(s')$ where s' is the state you get to after taking action a in state s
- ☐ $y = \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ☒ $y = R(s) + \gamma \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ☐ $y = R(s)$

 **Correct**

3.

1 / 1 point

You have reached the final practice quiz of this class! What does that mean? (Please check all the answers, because all of them are correct!)

- ☒ Andrew sends his heartfelt congratulations to you!

 **Correct**

- ☒ What an accomplishment -- you made it!

 **Correct**

- ☒ You deserve to celebrate!

 **Correct**

- ☒ The DeepLearning.AI and Stanford Online teams would like to give you a round of applause!

 **Correct**